

## 4. K Nearest Neighbors

---

- One of the easiest machine learning algorithms to implement
- Classification algorithm that makes a prediction based on a defined number of nearest neighbors
  - This is the only hyperparameter
  - Well, only one that's mandatory
- Doesn't require any "learning" - only a simple distance calculation
- Basic idea: A sample is classified by a majority vote of its nearest neighbors
- The algorithm classifies an instance based on the majority votes from k nearest instances
- Gotchas:
  - It's a distance-based algorithm, so it's a good idea to standardize your data before training

## Math behind

- KNN boils down to a single formula for distance
  - You can choose from many distance metrics, such as Cosine or Euclidean distance
  - We'll use Euclidean distance
  - Here's the formula:
- 
- Or in a more compact way:
- 
- It's a simple straight line distance that's easy to implement in Python

## Implementation

- The KNN class was written with having Scikit-Learn API syntax in mind

- The value for k is set by default to 3, and is stored to the constructor (`__init__()`) alongside `X_train` and `y_train`, initially set to `None`
- The `_euclidean_distance(p, q)` is a private method which implements the formula from above
- The `fit(X, y)` method does basically nothing - it just stores training data to the constructor
- The `predict(X)` method calculates distances between every row in X and every row in `X_train`. The distances are then sorted and only top k are kept. The classification is then made by calculating statistical mode

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from scipy import stats
from matplotlib import rcParams
rcParams['axes.spines.top'] = False
rcParams['axes.spines.right'] = False
```

In [2]:

```
class KNN:
    """
    A class which implement k Nearest Neighbors algorithm from scratch.
    """

    def __init__(self, k=3):
        self.k = k
        self.X_train = None
        self.y_train = None

    @staticmethod
    def _euclidean_distance(p, q):
        """
        Private method, calculates euclidean distance between two vectors.

        :param p: np.array, first vector
        :param q: np.array, second vector
        :return: float, distance
        """

        return np.sqrt(np.sum((p - q) ** 2))

    def fit(self, X, y):
        """
        Trains the model.
        No training is required for KNN, so `fit(X, y)` saves both parameters
        to the constructor.
        """
```

```

:param X: pd.DataFrame, features
:param y: pd.Series, target
:return: None
"""

self.X_train = X
self.y_train = y

def predict(self, X):
    """
    Predicts the class labels based on nearest neighbors.

    :param X: pd.DataFrame, features
    :return: np.array, predicted class labels
    """

    predictions = []
    for p in X:
        euc_distances = [self._euclidean_distance(p, q) for q in self.X_train]
        sorted_k = np.argsort(euc_distances)[:self.k]
        k_nearest = [self.y_train[y] for y in sorted_k]
        predictions.append(stats.mode(k_nearest)[0][0])

    return np.array(predictions)

```

## Testing

- We'll use the *breast cancer* dataset from Scikit-Learn:

In [3]:

```

from sklearn.datasets import load_breast_cancer

data = load_breast_cancer()
X = data.data
y = data.target

```

- The below code cell applies train/test split to the dataset:

In [4]:

```

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

```

- You can now initialize and "train" the model, and afterwards make predictions:

```
In [5]:  
model = KNN()  
model.fit(X_train, y_train)  
preds = model.predict(X_test)  
In [6]:  
preds  
  
In [7]:  
y_test
```

- As we can see, the arrays are quite similar but they have some differences
- Let's calculate accuracy to evaluate the model:

```
In [8]:  
from sklearn.metrics import accuracy_score  
  
accuracy_score(y_test, preds)
```

- The overall performance seems to be good out of the box
- Let's tweak it by finding the optimal K value

## K Optimization

- K value of 3 maybe isn't the best choice
- We'll train and evaluate the model for every odd number from 1 to 15
- The results are stored to evals:

```
In [9]:  
evals = []  
  
for k in range(1, 16, 2):  
    model = KNN(k=k)  
    model.fit(X_train, y_train)  
    preds = model.predict(X_test)  
    accuracy = accuracy_score(y_test, preds)  
    evals.append({'K': k, 'accuracy': accuracy})
```

- Let's visually inspect the best K:

```
In [10]:  
evals = pd.DataFrame(evals)  
best_k = evals.sort_values(by='accuracy', ascending=False).iloc[0]
```

```

plt.figure(figsize=(16, 8))
plt.plot(evals['k'], evals['accuracy'], lw=3, c='#087E8B')
plt.scatter(best_k['k'], best_k['accuracy'], s=200, c='#087E8B')
plt.title("K Parameter Optimization, Optimal k = " + str(int(best_k['k'])), size=20)
plt.xlabel('K', size=14)
plt.ylabel('Accuracy', size=14)
plt.savefig('images/004_KNN_k_optimization.png', dpi=300, bbox_inches='tight')
plt.show()

```

- The K value of 11 seems to work the best with our dataset
- You can now retrain the model with having this in mind
  - `model = KNN(k=11)`
- Let's compare performance with Scikit-Learn model next

## Comparison with Scikit-Learn

- We want to know if our model is good, so let's compare it with `KNeighborsClassifier` model from Scikit-Learn

In [11]:

```

from sklearn.neighbors import KNeighborsClassifier

knn_model = KNeighborsClassifier()
knn_model.fit(X_train, y_train)
knn_preds = knn_model.predict(X_test)

```

In [12]:

```
accuracy_score(y_test, knn_preds)
```

- They are more or less the same, but we didn't do K optimization for Scikit-Learn model