

Pandas tutorial 5: Scatter plot with pandas and matplotlib

Scatter plots are frequently used in data science and machine learning projects. In this pandas tutorial, I'll show you two simple methods to plot one. Both solutions will be equally useful and quick:

- one will be using **pandas** (more precisely: `pandas.plot.scatter()`)
- the other one using **matplotlib** (`matplotlib.pyplot.scatter()`)

Let's see them — and as usual: I'll guide you through step by step.

Note: If you don't know anything about pandas (or Python), you might want to start here:

1. *Python libraries and packages for Data Scientists*
2. [*Learn Python from Scratch*](#)
3. [*Pandas Tutorial 1 \(Basics\)*](#)
4. *Pandas Tutorial 2 (Aggregation and grouping)*
5. *Pandas Tutorial 3 (Data Formatting)*
6. *Pandas Tutorial 4 (Plotting in pandas: Bar Chart, Line Chart, Histogram)*

Download the code base!

This is a hands-on tutorial, so it's best if you do the coding part with me


What is a scatter plot? And what is it good for?

Scatter plots are used to visualize the relationship between two (or sometimes three) variables in a data set. The idea is simple:

- you take a data point,
- you take two of its variables,
- the y-axis shows the value of the first variable,
- the x-axis shows the value of the second variable

Following this concept, you display each and every datapoint in your dataset.

You'll get something like this:




Boom! **This is a scatter plot.** At least, the easiest (and most common) example of it.

This particular scatter plot shows the relationship between the height and weight of people from a random sample. Again:

- y-axis shows the height
- x-axis shows the weight
- and each blue dot represents a person in this dataset

So, for instance, this person's (*highlighted with red*) weight and height is 66.5 kg and 169 cm.




How to read a scatter plot?

Scatter plots play an important role in data science – **especially in building/prototyping machine learning models**. Looking at the chart above, you can immediately tell that there's a strong correlation between weight and height, right? As we discussed in my linear regression work, you can even fit a


trend line (a.k.a. regression line) to this data set and try to describe this relationship with a mathematical formula.

Something like this:




This above is called a **positive correlation**. The greater is the height value, the greater is the expected weight value, too. (*Of course, this is a generalization of the data set. There are always exceptions and outliers!*)

But it's also possible that you'll get a **negative correlation**:



And in real-life data science projects, you'll see **no correlation** often, too:



Anyway: if you see a sign of positive or negative correlation between two variables in a data science project, that's a good indicator that you found something interesting — something that's worth digging deeper into. Well, in 99% of cases it will turn out to be either a triviality, or a coincidence. But in the remaining 1%, *you might find gold!*

Okay, I hope I set your expectations about scatter plots high enough.

It's time to see how to create one in Python!

Scatter plot in pandas and matplotlib

As I mentioned before, I'll show you two ways to create your scatter plot.

You'll see here the Python code for:

- a **pandas scatter plot**
and
- a **matplotlib scatter plot**

The two solutions are fairly similar, the whole process is ~90% the same...

The only difference is in the last few lines of code.

Note: By the way, I prefer the matplotlib solution because I find it a bit more transparent.

I'll guide you through these 4 steps:

1. Importing pandas, numpy and matplotlib
2. Getting the data
3. Preparing the data
4. Plotting a scatter plot

Step #1: Import pandas, numpy and matplotlib!

Just as we have done in the histogram work, as a first step, you'll have to import the libraries you'll use. And you'll also have to make a small tweak in your Jupyter environment.

```
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
  
%matplotlib inline
```

The first two lines will import `pandas` and `numpy`.

The third line will import the `pyplot` from `matplotlib` — also, we will refer to it as `plt`.

And `%matplotlib inline` sets your environment so you can directly plot charts into your Jupyter Notebook!

Great!

```
In [1]: import numpy as np  
       import pandas as pd  
       import matplotlib.pyplot as plt  
       %matplotlib inline
```

Step #2: Get the data!

Well, in real data science projects, getting the data would be a bit harder. You should read .csv files or SQL tables into your Python environment. But this tutorial's focus is not on learning that — so you can take the lazy way and use the dataset I'll provide for you here.

Note: What's in the data? This is the modified version of the dataset that we used in the pandas histogram work — the heights and weights of our hypothetical gym's members.

```
np.random.seed(0)
mu = 170 #mean
sigma = 6 #stddev
sample = 100
height = np.random.normal(mu, sigma, sample)

weight = (height-100) * np.random.uniform(0.75, 1.25, 100)
```

This is a random generator, by the way, that generates 100 height and 100 weight values — in numpy array format. By using the `np.random.seed(0)` line, we also made sure you'll be able to work with the exact same data points that I do in this article.

```
In [4]: height
```

```
Out[4]: array([180.58431408, 172.40094325, 175.8724279 , 183.4453592 ,  
181.20534794, 164.13633272, 175.70053051, 169.09185675,  
169.38068689, 172.46359101, 170.86426143, 178.72564104,  
174.56622635, 170.7300501 , 172.6631794 , 172.00204596,  
178.96447444, 168.76905042, 171.87840621, 164.87542556,  
154.6820611 , 173.92171157, 175.18661719, 165.54700988,  
183.61852774, 161.27380595, 170.2745511 , 168.8768969 ,  
179 19667529 178 81615262 170 92968455 172 26897512
```

```
In [5]: weight
```

```
Out[5]: array([77.51626973, 76.25242778, 57.63243821, 75.16652885, 87.70882186,  
57.40450651, 80.1674458 , 66.6311054 , 56.73515698, 65.1549975 ,  
73.34326729, 82.30264923, 77.33730302, 76.14800093, 78.18933299,  
69.53303948, 94.6210212 , 64.2152282 , 69.57344273, 77.58852279,  
63.05372031, 81.45760811, 60.1578232 , 79.29492539, 92.57579877,  
76.55693329, 57.95711959, 81.55458715, 65.83195656, 83.37013275,  
57 58851958 84 84407415 74 61094889 60 12312355 64 76081868
```

Note: For now, you don't have to know line by line what's going on here. (I'll write a separate article about how `numpy.random` works.)

In the next step, we will push these data sets into pandas dataframes.

Step #3: Prepare the data!

Again, preparing, cleaning and formatting the data is a painful and time consuming process in real-life data science projects. But in this tutorial, we are lucky, everything is prepared – the data is clean – so you can push your `height` and `weight` data sets directly into a pandas dataframe (called `gym`) by running this one line of code:

```
gym = pd.DataFrame({'height': height, 'weight': weight})
```

Note: If you want to experience the complexity of a true-to-life data science project, go and check out my 6-week course: [The Junior Data Scientist's First Month!](#)

Your `gym` dataframe should look like this.

In [9]: `gym`

Out[9]:

	height	weight
0	180.584314	77.516270
1	172.400943	76.252428
2	175.872428	57.632438
3	183.445359	75.166529
4	181.205348	87.708822
...
95	174.239439	86.527865
96	170.063000	74.581428
97	180.715223	95.754513
98	170.761473	62.749223
99	172.411936	83.203010

100 rows × 2 columns

Perfect: ready for putting it on a scatter plot!

Step #4a: Pandas scatter plot

Okay, all set, we have the `gym` dataframe. Let's create a **pandas scatter plot!**

Now, this is only one line of code and it's pretty similar to what we had for bar charts, line charts and histograms in pandas...

It starts with: `gym.plot ...`and then you simply have to define the chart type that you want to plot, which is `scatter()`. But when plotting a scatter plot in pandas, you'll always have to specify the `x` and `y` values as parameters, too.
(This could seem unusual because for bar and line charts, you didn't have to do anything similar to this.)


So the final line of code will be:

```
gym.plot.scatter(x = 'weight', y = 'height')
```

The `x` and `y` values – by definition – have to come from the `gym` dataframe, so you have to refer to the column names: `'weight'` and `'height'`!

A quick comment: Watch out for all the apostrophes! I know from my live workshops that the syntax might seem tricky at first. But you'll get used to it after your 5th or 6th scatter plot, I promise!

That's it! You have plotted a scatter plot in pandas!



Step #4b: Matplotlib scatter plot

Here's an alternative solution for the last step. In this one, we will use the `matplotlib` library instead of `pandas`. (*Although, I have to mention here that the pandas solution I showed you is actually built on matplotlib's code.*)

In my opinion, this solution is a bit more elegant. But from a technical standpoint — and for results — both solutions are equally great.

Anyway, type and run these three lines:

```
x = gym.weight  
y = gym.height
```

```
plt.scatter(x,y)
```

- `x = gym.weight`

This line defines what values will be displayed on the x-axis of the scatter plot. It's the `weight` column again from the `gym` dataset. (Note: *This is in pandas Series format... But in this specific case, I could have added the original numpy array, too.*)

- `y = gym.height`

On the y-axis we want to display the `gym.height` values. (*This is in pandas Series format, too!*)


- `plt.scatter(x,y)`

And then this line does the plotting. Remember, you set `plt` at the very beginning of our Jupyter notebook (`import matplotlib.pyplot as plt`) — and so `plt` refers to `matplotlib.pyplot`! And the `x` and `y` values are parameters that have been defined in the previous two lines.

```
In [8]: x = gym.weight  
y = gym.height  
plt.scatter(x,y)
```

Again: this is slightly different (and in my opinion slightly nicer) syntax than with pandas.

But the result is exactly the same.



Conclusion

This is how you make a scatter plot in pandas and/or in matplotlib. I think it's fairly easy and I hope you think the same. If you haven't done so yet, check out my Python histogram tutorial, too! If you have any questions, leave a comment below!